

1 Topological Artist Model

To guide the implementation of structure preserving building block components, we develop a mathematical formalism of visualization that specifies how these components preserve *continuity* and *equivariance*. Inspired by the somewhat analogous component in Matplotlib[1], we call the transformation from data space to graphic that these building block components implement the *artist*.

$$\mathcal{A} : \mathcal{E} \rightarrow \mathcal{H} \quad (1)$$

The *artist* \mathcal{A} is a map from the data \mathcal{E} to graphic \mathcal{H} fiber bundles. To explain how the *artist* is a structure preserving map from data to graphic, we first describe how we model data (1.1) and graphics (1.2) as topological structures that encapsulate component types and continuity. We then discuss the maps from graphic to data (1.2.2, data components to visual components (1.3.2), and visual components into graphic (1.3.3) that make up the artist.

1.1 Data Space E

Building on Butler’s proposal of using fiber bundles as a common data representation structure for visualization data[2, 3], a fiber bundle is a tuple (E, K, π, F) defined by the projection map π

$$F \hookrightarrow E \xrightarrow{\pi} K \quad (2)$$

that binds the components of the data in F to the continuity represented in K . The fiber bundle models the properties of data component types F (1.1.1), the continuity of records K (1.1.3), the collections of records τ (1.1.4), and the space E of all possible datasets with these components and continuity.

By definition fiber bundles are locally trivial[4, 5], meaning that over a localized neighborhood we can dispense with extra structure on E and focus on the components and continuity. We use fiber bundles as the data model because they are inclusive enough to express all the types of data described in section ??.

1.1.1 Variables in Fiber Space F

To formalize the structure of the data components, we use notation introduced by Spivak [6] that binds the components of the fiber to variable names. This allows us to describe the components in a schema like way. Spivak constructs a set \mathbb{U} that is the disjoint union of all possible objects of types $\{T_0, \dots, T_m\} \in \mathbf{DT}$, where \mathbf{DT} are the data types of the variables in the dataset. He then defines the single variable set \mathbb{U}_σ

$$\begin{array}{ccc} \mathbb{U}_\sigma & \longrightarrow & \mathbb{U} \\ \pi_\sigma \downarrow & & \downarrow \pi \\ C & \xrightarrow{\sigma} & \mathbf{DT} \end{array} \quad (3)$$

which is \mathbb{U} restricted to objects of type T bound to variable name c . The \mathbb{U}_σ lookup is by name to specify that every component is distinct, since multiple components can have the same type T . Given σ , the fiber for a one variable dataset is

$$F = \mathbb{U}_{\sigma(c)} = \mathbb{U}_T \quad (4)$$

where σ is the schema binding variable name c to its datatype T . A dataset with multiple variables has a fiber that is the cartesian cross product of \mathbb{U}_σ applied to all the columns:

$$F = \mathbb{U}_{\sigma(c_1)} \times \dots \times \mathbb{U}_{\sigma(c_i)} \dots \times \mathbb{U}_{\sigma(c_n)} \quad (5)$$

which is equivalent to

$$F = F_0 \times \dots \times F_i \times \dots \times F_n \quad (6)$$

18 which allows us to decouple F into components F_i .




Figure 1: These two datasets have the same continuity, but different components and therefore different fibers. The 2D fiber $F = \mathbb{R} \times \mathbb{R}$ encodes a dataset consisting of *time* and *temperature* components. One dimension of the fiber encodes the range of possible values for the time component of the dataset, which is a subset of the \mathbb{R} , while the other dimension encodes the range of possible values \mathbb{R} for the temperature component. The 3D fiber encodes a dataset consisting of *temperature*, *latitude*, and *longitude* components.

For example, the records in the 2D fiber in [Figure 1](#) are a pair of times and °K temperature measurements taken at those times. Time is a positive number of type `datetime` which can be resolved to floats $\mathbb{U}_{\text{datetime}} = \mathbb{R}$. Temperature values are real positive numbers $\mathbb{U}_{\text{float}} = \mathbb{R}^+$. The fiber is

$$F = \mathbb{R} \times \mathbb{R}^+$$

where the first component F_0 is the set of values specified by $(c = \text{time}, T = \text{datetime}, \mathbb{U}_\sigma = \mathbb{R})$ and F_1 is specified by $(c = \text{temperature}, T = \text{float}, \mathbb{U}_\sigma = \mathbb{R})$ and is the set of values $\mathbb{U}_\sigma = \mathbb{R}$. In the 3D fiber in [??](#), time is replaced with location. This location variable is of type `point` and has two components *latitude* and *longitude* $\{(lat, lon) \in \mathbb{R}^2 \mid -90 \leq lat \leq 90, 0 \leq lon \leq 360\}$. The fiber for this dataset is

$$F = \mathbb{R} \times \mathbb{R}^2 = \mathbb{R} \times \mathbb{R} \times \mathbb{R}$$

19 where the dimensionality of the fiber does not change, but the components of the fiber
20 can be coupled. For example, *location* can either be specified as $(c = \text{location}, T =$

21 `point`, $\mathbb{U}_\sigma = \mathbb{R}^2$) or $(c = \text{latitude}, T = \text{float}, \mathbb{U}_\sigma = \mathbb{R})$ and $(c = \text{longitude}, T =$
 22 `float`, $\mathbb{U}_\sigma = \mathbb{R})$.

23 As illustrated in figure 1, Spivak’s framework provides a consistent way to describe
 24 potentially complex components of the input data.

25 1.1.2 Measurement Scales: Monoid Actions

26 Implementing expressive visual encodings requires formally describing the structure on the
 27 components of the fiber, which we define by the actions of a monoid on the component. In
 28 doing so, we specify the properties of the component that must be preserved in a graphic
 29 representation. While structure on a set of values is often described algebraically as oper-
 30 ations or through the actions of a group, for example Steven’s scales [7], we generalize to
 31 monoids to support more component types. Monoids are also commonly found in functional
 32 programming because they specify compositions of transformations [8, 9].

A monoid [10] M is a set with an associative binary operator $*$: $M \times M \rightarrow M$. A
 monoid has an identity element $e \in M$ such that $e * a = a * e = a$ for all $a \in M$. As
 defined on a component of F , a left monoid action [11, 12] of M_i is a set F_i with an action
 $\bullet : M \times F_i \rightarrow F_i$ with the properties:

associativity for all $f, g \in M_i$ and $x \in F_i$, $f \bullet (g \bullet x) = (f * g) \bullet x$


identity for all $x \in F_i$, $e \in M_i$, $e \bullet x = x$

As with the fiber F the total monoid space M is the cartesian product


$$M = M_0 \times \dots \times M_i \times \dots \times M_n \quad (7)$$

33 of each monoid M_i on F_i . The monoid is also added to the specification of the fiber
 34 $(c_i, T_i, \mathbb{U}_\sigma M_i)$

35 Steven’s described the measurement scales [7, 13] in terms of the monoid actions
 36 on the measurements: nominal data is permutable, ordinal data is monotonic, in-
 37 terval data is translatable, and ratio data is scalable [14]. For example, lets say
 38 one component of the data is weather states, such that the fiber component is $F_i =$
 39 $\{\text{sun}, \text{clouds}, \text{rain}, \text{lightning}, \text{storm}\}$. If the data is ordinal, that means some ordering
 40 action is applied to the set.



(a) The identity action \bullet maps the weather states to themselves.



(b) The action \bullet is a monotonic map, meaning that it preserves the relative ordering despite mapping multiple elements to the same element.

41 I think there's something about associativity here, but I'm not sure what In Figure 2b,
 42 the action \bullet acts on the weather states such that the states that they are mapped into have
 43 the same relative ordering as the original states. Unlike group actions, the action \bullet is not
 44 invertable and therefore a monoid action. We generalize to monoids specifically for the case
 45 of partially ordered scales are common in indicator rankings [15].

46 1.1.3 Continuity of the Data K

47 The base space K is way to express how the records in E are connected to each other, for
 48 example if they are discrete points or if they lie in a 2D continous surface. Connectivity
 49 type is assumed in the choice of visualization, for example a line plot implies 1D continuous
 50 data, but an explicit representation allows for verifying that the topology of the graphic
 51 representation is equivalent to the topology of the data.



Figure 3: The topological base space K encodes the continuity of the data space, for example if the data is discrete points or lies on a plane or a sphere

52 As illustrated in figure 3, K is akin to an indexing space into E that describes the
 53 structure of E . K can have any number of dimensions and can be continuous or discrete.




Figure 4: The base space E is divided into fiber segments F . The base space K acts as an index into the records in the fibers.

Formally K is the quotient space [16] of E meaning it is the finest space[17] such that every $k \in K$ has a corresponding fiber F_k [16]. In figure 4, E is a rectangle divided by vertical fibers F , so the minimal K for which there is always a mapping $\pi : E \rightarrow K$ is the closed interval $[0, 1]$. As with fibers and monoids, we can decompose the total space into components $\pi : E_i \rightarrow K$ where

$$\pi : E_1 \oplus \dots \oplus E_i \oplus \dots \oplus E_n \rightarrow K \quad (8)$$

54 which is a decomposition of F . The K remains the same because the connectivity of records
 55 does not change just because there are fewer elements in each record.




Figure 5: These two datasets have the same (time, temperature) fiber, but different continuities. The dataset on the left consists of discrete records, while the records in the dataset on the right sampled from a continuous space.

56 The datasets in figure 5 have the same fiber of (temperature, time). The dot represents
 57 a discrete base space K , meaning that every dataset encoded in the fiber bundle has discrete
 58 continuity. The line is a representation of a 1D continuity, meaning that every dataset in the
 59 fiber bundle is 1D continuous. By encoding this continuity in the model as K the data model
 60 now explicitly carries information about its structure such that the implicit assumptions of
 61 the visualization algorithms are now explicit. The explicit topology is a concise way of
 62 distinguishing visualizations that appear identical, for example heatmaps and images.

63 1.1.4 Data τ

While the projection function $\pi : E \rightarrow K$ ties together the base space K with the fiber F , a section $\tau : K \rightarrow E$ encodes a dataset. A section function takes as input location $k \in K$ and returns a record $r \in E$. For example, in the special case of a table [6], K is a set of row ids, F is the columns, and the section τ returns the record r at a given key in K . For any fiber bundle, there exists a map

$$\begin{array}{ccc} F & \hookrightarrow & E \\ & \pi \downarrow & \uparrow \tau \\ & K & \end{array} \quad (9)$$

such that $\pi(\tau(k)) = k$. The set of all global sections is denoted as $\Gamma(E)$. Assuming a trivial fiber bundle $E = K \times F$, the section is

$$\tau(k) = (k, (g_{F_0}(k), \dots, g_{F_n}(k))) \quad (10)$$

where $g : K \rightarrow F$ is the index function into the fiber. This formulation of the section also holds on locally trivial sections of a non-trivial fiber bundle. Because we can decompose the bundle and the fiber, we can decompose τ as

$$\tau = (\tau_0, \dots, \tau_i, \dots, \tau_n) \quad (11)$$

64 where each section τ_i is a variable or set of variables. This allows for accessing the data
 65 component wise in addition to accessing the data in terms of its location over K .




Figure 6: Fiber (time, temperature) with an interval K basespace. The sections $\tau^{(1)}$ and $\tau^{(2)}$ are constrained such that the time variable must be monotonic, which means each section is a timeseries of temperature values. They are included in the global set of sections $\tau^{(1)}, \tau^{(2)} \in \Gamma(E)$

66 In the example in figure 6, the fiber is $(time, temperature)$ as described in figure 1 and
 67 the base space is the interval K . The section $\tau^{(1)}$ resolves to a series of monotonically
 68 increasing in time records of (time, temperature) values. Section $\tau^{(2)}$ returns a different
 69 timeseries of (time, temperature) values. Both sections are included in the global set of
 70 sections $\tau^{(1)}, \tau^{(2)} \in \Gamma(E)$.

71 1.1.5 Applications to Data Containers

72 This model provides a common formalism for widely used data containers without sacrificing
 73 the semantic structure embedded in each container. For example, the section can be any
 74 instance of a univariate numpy array[18] that stores an image. This could be a section of a
 75 fiber bundle where K is a 2D continuous plane and the F is $(\mathbb{R}^3, \mathbb{R}, \mathbb{R})$ where \mathbb{R}^3 is color,
 76 and the other two components are the x and y positions of the sampled data in the image.
 77 This position information is already implicitly encoded in the array as the index and the
 78 resolution of the image being stored. Instead of an image, the numpy array could also store a
 79 2D discrete table. The fiber would not change, but the K would now be 0D discrete points.
 80 These different choices in topology indicate, for example, what sorts of interpolation would
 81 be appropriate when visualizing the data.

82 There are also many types of labeled containers that can richly be described in this
 83 framework because of the schema like structure of the fiber. For example, a pandas series
 84 which stores a labeled list, or a dataframe[19] which stores a relational table. A series could
 85 store the values of $\tau^{(1)}$ and a second series could be $\tau^{(2)}$. We could also fatten the fiber to
 86 hold two temperature series, such that a section would be an instance of a dataframe with
 87 a time column and two temperature columns. While the series and dataframe explicitly
 88 have a time index column, they are components in our model and the index is assumed to
 89 be data independent references such as hashvalues, virtual memory locations, or random
 90 number keys.

91 Where this model particularly shines are N dimensional labeled data structures. For
 92 example, an xarray[20] data that stores temperature field could have a K that is a continuous
 93 volume and the components would be the temperature and the time, latitude, and longitude

the measurements were sampled at. A section can also be an instance of a distributed data container, such as a dask array [21]. As with the other containers, K and F are defined in terms of the index and dtypes of the components of the array. Because our framework is defined in terms of the fiber, continuity, and sections, rather than the exact values of the data, our model does not need to know what the exact values are until the renderer needs to fill in the image.

1.2 Graphic Space H

We introduce a graphic bundle to hold the essential information necessary to render a graphical design constructed by the artist. As with the data, we can represent the target graphic as a section ρ of a bundle (H, S, π, D) . The graphic bundle H consists of a base S (1.2.1) that is a thickened form of K a fiber D (1.2.2) that is an idealized display space, and sections ρ (1.2.3) that encode a graphic where the visual characteristics are fully specified.

1.2.1 Idealized Display D

To fully specify the visual characteristics of the image, we construct a fiber D that is an infinite resolution version of the target space. Typically H is trivial and therefore sections can be thought of as mappings into D . In this work, we assume a 2D opaque image $D = \mathbb{R}^5$ with elements

$$(x, y, r, g, b) \in D \quad (12)$$

such that a rendered graphic only consists of 2D position and color. To support overplotting and transparency, the fiber could be $D = \mathbb{R}^7$ such that $(x, y, z, r, g, b, a) \in D$ specifies the target display. By abstracting the target display space as D , the model can support different targets, such as a 2D screen or 3D printer.

1.2.2 Continuity of the Graphic S

Just as the K encodes the connectivity of the records in the data, we propose an equivalent S that encodes the connectivity of the rendered elements of the graphic. For example, consider a S that is mapped to the region of a 2D display space that represents K . For some visualizations, K may be lower dimension than S . For example, a point that is 0D in K cannot be represented on screen unless it is thickened to 2D to encode the connectivity of the pixels that visually represent the point. This thickening is often not necessary when the dimensionality of K matches the dimensionality of the target space, for example if K is 2D and the display is a 2D screen. We introduce S to thicken K in a way which preserves the structure of K .

Formally, we require that K be a deformation retract[22] of S so that K and S have the same homotopy. The surjective map $\xi : S \rightarrow K$

$$\begin{array}{ccc} E & & H \\ \pi \downarrow & & \pi \downarrow \\ K & \xleftarrow{\xi} & S \end{array} \quad (13)$$

goes from region $s \in S_k$ to its associated point s . This means that if $\xi(s) = k$, the record at k is copied over the region s such that $\tau(k) = \xi^* \tau(s)$ where $\xi^* \tau(s)$ is τ pulled back over S .




Figure 7: The scatter and line graphic base spaces have one more dimension of continuity than K so that S can encode physical aspects of the glyph, such as shape (a circle) or thickness. The image has the same dimension in S as in K .

When K is discrete points and the graphic is a scatter plot, each point $k \in K$ corresponds to a 2D disk S_k as shown in figure 7. In the case of 1D continuous data and a line plot, the region β over a point α_i specifies the thickness of the line in S for the corresponding τ on k . The image has the same dimensions in data space and graphic space such that no extra dimensions are needed in S .

The mapping function ξ provides a way to identify the part of the visual transformation that is specific to the connectivity of the data rather than the values; for example it is common to flip a matrix when displaying an image. The ξ mapping is also used by interactive visualization components to look up the data associated with a region on screen. One example is to fill in details in a hover tooltip, another is to convert region selection (such as zooming) on S to a query on the data to access the corresponding record components on K .

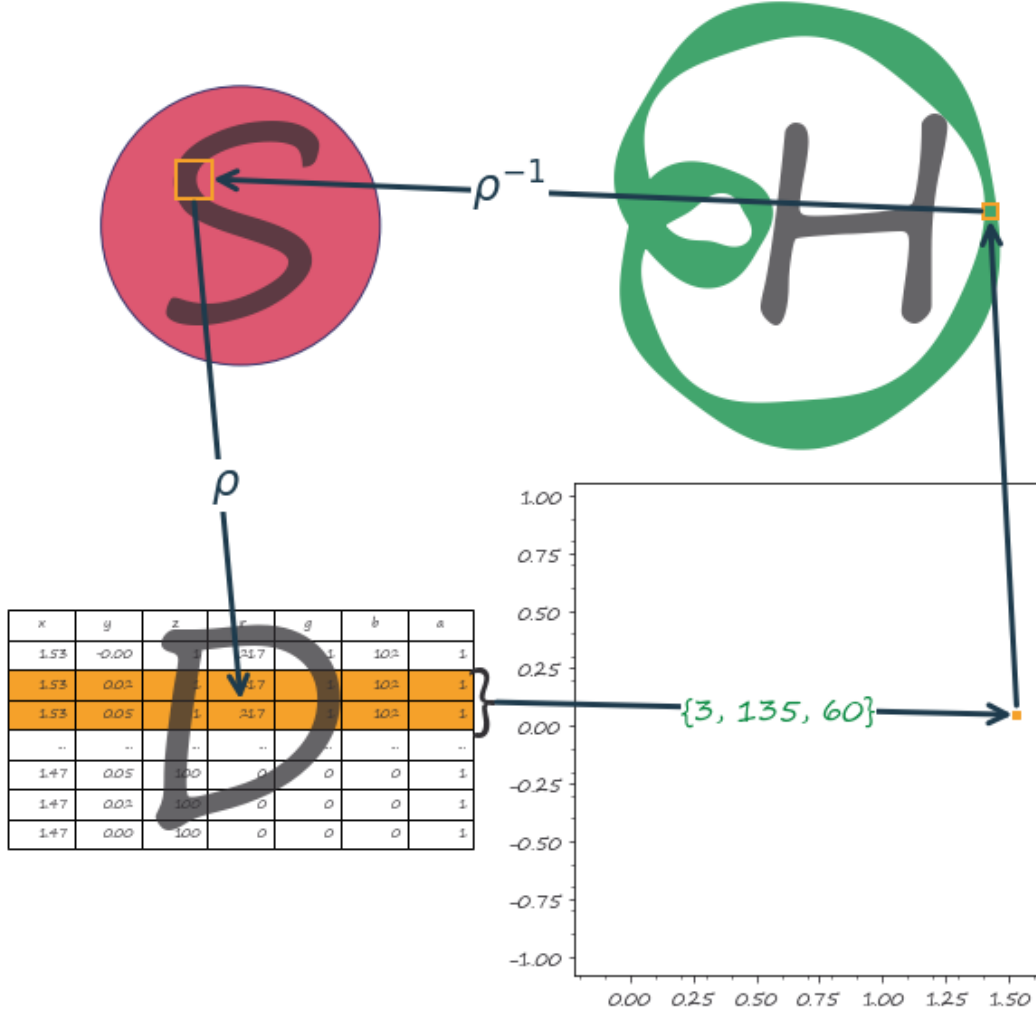


Figure 8: To render a graphic, a pixel p is selected in the display space, which is defined in the same coordinates as the x and y components in D . The inverse mapping $\rho_{xy}^{-1}(p)$ returns a region $S_p \subset S$. $\rho(S_p)$ returns the list of elements $(x, y, r, g, b) \in D$ that lie over S_p . The integral over the (r, g, b) elements is the color of the pixel.

136 This section describes how we go from a graphic in an idealized prerender space to a rendered
 137 image, where the graphic is the section $\rho : S \rightarrow H$. It is sufficient to sketch out how an
 138 arbitrary pixel would be rendered, where a pixel p in a real display corresponds to a region
 139 S_p in the idealized display. To determine the color of the pixel, we aggregate the color values
 140 over the region via integration.

For a 2D screen, the pixel is defined as a region $p = [y_{top}, y_{bottom}, x_{right}, x_{left}]$ of the rendered graphic. Since the x and y in p are in the same coordinate system as the x and y components of D the inverse map of the bounding box $S_p = \rho_{xy}^{-1}(p)$ is a region $S_p \subset S$. To compute the color, we integrate on S_p

$$r_p = \iint_{S_p} \rho_r(s) ds^2 \quad (14)$$

$$g_p = \iint_{S_p} \rho_g(s) ds^2 \quad (15)$$

$$b_p = \iint_{S_p} \rho_b(s) ds^2 \quad (16)$$

As shown in figure 8, a pixel p in the output space is selected and inverse mapped into the corresponding region $S_p \subset S$. This triggers a lookup of the ρ over the region S_p , which yields the set of elements in D that specify the (r, g, b) values corresponding to the region p . The color of the pixel is then obtained by taking the integral of $\rho_{rgb}(S_p)$.

In general, ρ is an abstraction of rendering. In very broad strokes ρ can be a specification such as PDF[23], SVG[24], or an OpenGL scene graph[25]. Alternatively, ρ can be a rendering engine such as cairo[26] or AGG[**AntiGrainGeometry**]. Implementation of ρ is out of scope for this work,

1.3 Artist

We propose that the transformation from data to visual representation can be described as a structure preserving map from one topological space to another. We name this map the artist as that is the analogous part of the Matplotlib[1] architecture that builds visual elements. The topological artist A is a monoid equivariant sheaf map from the sheaf on a data bundle E which is $\mathcal{O}(E)$ to the sheaf on the graphic bundle H , $\mathcal{O}(H)$.

$$A : \mathcal{O}(E) \rightarrow \mathcal{O}(H) \quad (17)$$

Sheafs are a mathematical object with restriction maps that define how to glue τ over local neighborhoods $U \subseteq K$, discussed in section ??, such that the A maps are consistent over continuous regions of K . While A can usually construct graphical elements solely with the data in τ , some visualizations, such as line, may also need some finite number n of derivatives, which is captured by the jet bundle \mathcal{J}^n [27, 28] with $\mathcal{J}^0(E) = E$. In this work, we at most need $\mathcal{J}^2(E)$ which is the value at τ and its first and second derivatives; therefore the artist takes as input the jet bundle $E' = \mathcal{J}^2(E)$.

Specifically, A is the equivariant map from E' to a specific graphic $\rho \in \Gamma(H)$

$$\begin{array}{ccccc} E' & \xrightarrow{\nu} & V & \xleftarrow{\xi^*} & \xi^* V & \xrightarrow{Q} & H \\ & \searrow \pi & \downarrow \pi & & \downarrow \xi^* \pi & & \swarrow \pi \\ & & K & \xleftarrow{\xi} & S & & \end{array} \quad (18)$$

where the input can be point wise $\tau(k) \mid k \in K$. The encoders $\nu : E' \rightarrow V$ convert the data components to visual components(1.2.2). The continuity map $\xi : S \rightarrow K$ then

pulls back the visual bundle V over $S(1.3.2)$. Then the assembly function $Q : \xi^*V \rightarrow H$ composites the fiber components of ξ^*V into a graphic in $H(1.3.3)$. This functional decomposition of the visualization artist facilitates building reusable components at each stage of the transformation because the equivariance constraints are defined on ν , Q , and ξ .

1.3.1 Visual Fiber Bundle V

We introduce a visual bundle V to store the visual representations the artist needs to assemble into a graphic. The visual bundle (V, K, π, P) has section $\mu : V \rightarrow K$ that resolves to a visual variable in the fiber P . The visual bundle V is the latent space of possible parameters of a visualization type, such as a scatter or line plot. We define P in terms of the parameters of a visualization libraries compositing functions; for example table 1 is a sample of the fiber space for Matplotlib [29].

ν_i	μ_i	$\text{codomain}(\nu_i) \subset P_i$
position	x, y, z, theta, r	\mathbb{R}
size	linewidth, markersize	\mathbb{R}^+
shape	markerstyle	$\{f_0, \dots, f_n\}$
color	color, facecolor, markerfacecolor, edgecolor	\mathbb{R}^4
texture	hatch	\mathbb{N}^{10}
	linestyle	$(\mathbb{R}, \mathbb{R}^{+n, n\%2=0})$

Table 1: Some possible components of the fiber P for a visualization function implemented in Matplotlib

A section μ is a tuple of visual values that specifies the visual characteristics of a part of the graphic. For example, given a fiber of $\{xpos, ypos, color\}$ one possible section could be $\{.5, .5, (255, 20, 147)\}$. The $\text{codomain}(\nu_i)$ determines the monoid actions on P_i . These fiber components are implicit in the library, by making them explicit as components of the fiber we can build consistent definitions and expectations of how these parameters behave.

180 **1.3.2 Visual Encoders ν**




Figure 9: In this artist, ν maps the strings to the emojis. This ν is equivariant because the monoid actions (which are represented by the colored arrows) are the same on both the τ input and μ output sets.

As introduced in section ??, there are many ways to visually represent data components. We define the visual transformers ν

$$\{\nu_0, \dots, \nu_n\} : \{\tau_0, \dots, \tau_n\} \mapsto \{\mu_0, \dots, \mu_n\} \quad (19)$$

181 as the set of equivariant maps $\nu_i : \tau_i \mapsto \mu_i$. Given M_i is the monoid action on E_i and that
 182 there is a monoid M'_i on V_i , then there is a monoid homomorphism from $\varphi : M_i \rightarrow M'_i$
 183 that ν must preserve. As mentioned in section 1.1.2, we choose monoid actions as the basis
 184 for equivariance because they define the structure on the fiber components.

A validly constructed ν is one where the diagram of the monoid transform m commutes such that

$$\begin{array}{ccc} E_i & \xrightarrow{\nu_i} & V_i \\ m_r \downarrow & & \downarrow m_v \\ E_i & \xrightarrow{\nu_i} & V_i \end{array} \quad (20)$$

In general, the data fiber F_i cannot be assumed to be of the same type as the visual fiber P_i and the actions of M on F_i cannot be assumed to be the same as the actions of M' on P_i ; therefore an equivariant ν_i must satisfy the constraint

$$\nu_i(m_r(E_i)) = \varphi(m_r)(\nu_i(E_i)) \quad (21)$$

185 such that φ maps a monoid action on data to a monoid action on visual elements. However,
 186 we can construct a monoid action of M on P_i that is compatible with a monoid action of
 187 M on F_i . We can compose the monoid actions on the visual fiber $M' \times P_i \rightarrow P_i$ with the
 188 homomorphism φ that takes M to M' . This allows us to define a monoid action on P of M

189 that is $(m, v) \rightarrow \varphi(m) \bullet v$. Therefore, without a loss of generality, we can assume that an
 190 action of M acts on F_i and on P_i compatibly such that φ is the identity function.

On example of an equivariant ν is illustrated in figure 9, which is a mapping from **Strings** to symbols. The data is an example of a Steven's nominal measurement set, which is defined as having on it permutation group actions

$$\text{if } r_1 \neq r_2 \text{ then } \nu(r_1) \neq \nu(r_2) \quad (22)$$

such that shuffling the words must have an equivalent shuffle of the symbols they are mapped to. This is illustrated in the identical actions, represented by the colored arrows, on the words and emojis. To preserve ordinal and partial order monoid actions, ν must be a monotonic function such that given $r_1, r_2 \in E_i$,

$$\text{if } r_1 \leq r_2 \text{ then } \nu(r_1) \leq \nu(r_2) \quad (23)$$

the visual encodings must also have some sort of ordering. For interval scale data, ν is equivariant under translation monoid actions if

$$\nu(x + c) = \nu(x) + c \quad (24)$$

while for ratio data, there must be equivalent scaling[14]

$$\nu(xc) = \nu(x) * c \quad (25)$$

We therefore can test if a ν is equivariant by testing the actions under which is must commute. For example, we define a transform $\nu_i(x) = .5$ on interval data. This means it must commute under translation, for example $t(x) = x + 2$. Testing this constraint

$$\nu(t(r + 2)) \stackrel{?}{=} \nu(r) + 2 \quad (26)$$

$$.5 \neq .5 + 2 \quad (27)$$

191 we find that the ν defined here does not commute and is therefore invalid. The constraints
 192 on ν can be embedded into our artist such that the ν functions can test for equivariance
 193 and also provide guidance on constructing new ν functions.




Figure 10: ν_i functions convert data τ_i to visual characteristics μ_i , then Q assembles μ_i into a graphic ρ such that there is a map ξ preserving the continuity of the data. ρ applied to a region of connected components S_j generates a part of a graphic, for example the point graphical mark.

195 As shown in figure 10, the assembly function Q combines the fiber F_i wise ν transforms into
 196 a graphic in H . Together, ν and Q are a map-reduce operation: map the data into their
 197 visual encodings, reduce the encodings into a graphic. As with ν the constraint on Q is
 198 that for every monoid action on the input μ there is corresponding monoid action on the
 199 output ρ .

While ρ generates the entire graphic, we will restrict the discussion of Q to generation of sections of a glyph. We formally describe a glyph as Q applied to the regions k that map back to a set of path connected components $J \subset K$ as input:

$$J = \{j \in K \text{ s. t. } \exists \gamma \text{ s.t. } \gamma(0) = k \text{ and } \gamma(1) = j\} \quad (28)$$

where the path[30] γ from k to j is a continuous function from the interval $[0,1]$. We define the glyph as the graphic generated by $Q(S_j)$

$$H \xrightleftharpoons[\rho(S_j)]{} S_j \xrightleftharpoons[\xi^{-1}(J)]{\xi(s)} J_k \quad (29)$$

200 such that for every glyph there is at least one corresponding region on K . This is in keeping
 201 with the definition of glyph as any differentiable element put forth by Ziemkiewicz and
 202 Kosara[31]. The primitive point, line, and area marks[32, 33] are specially cased glyphs.

203 It is on sections of these glyphs that we define the equivariant map as $Q : \mu \mapsto \rho$ and an
 204 action on the subset of graphics $Q(\Gamma(V)) \in \Gamma(H)$ that Q can generate. We then define the
 205 constraint on Q such that if Q is applied to μ, μ' that generate the same ρ then the output
 206 of both sections acted on by the same monoid m must be the same. While it may seem
 207 intuitive that visualizations that generate the same glyph should consistently generate the

208 same glyph given the same input, we formalize this constraint such that it can be specified
 209 as part of the implementation of Q .

Lets call the visual representations of the components $\Gamma(V) = X$ and the graphic $Q(\Gamma(V)) = Y$. If for elements of the monoid $m \in M$ and for all $\mu, \mu' \in X$, we define the monoid action on X so that it is by definition equivariant

$$Q(\mu) = Q(\mu') \implies Q(m \circ \mu) = Q(m \circ \mu') \quad (30)$$

210 then a monoid action on Y can be defined as $m \circ \rho = \rho'$. The transformed graphic ρ' is
 211 equivariant to a transform on the visual bundle $\rho' = Q(m \circ \mu)$ on a section that $\mu \in Q^{-1}(\rho)$
 212 that must be part of generating ρ .




Figure 11: These two glyphs are generated by the same Q function. The monoid action m_i on edge thickness μ_i of the first glyph yields the thicker edge μ_i' in the second glyph.

213 The glyph in figure 11 has the following characteristics P specified by $(xpos, ypos, color, thickness)$
 214 such that one section is $\mu = (0, 0, 0, 1)$ and $Q(\mu) = \rho$ generates a piece of the thin hollow
 215 circle. The equivariance constraint on Q is that the action $m = (e, e, e, x + 2)$, where e is
 216 identity, translates μ to $\mu' = (e, e, e, 3)$. The corresponding action on ρ causes $Q(\mu')$ to be
 217 the thicker circle in figure 11.

218 1.3.4 Assembly Q

219 In this section we formulate the minimal Q that will generate distinguishable graphical
 220 marks: non-overlapping scatter points, a non-infinitely thin line, and an image.




Figure 12: The data is discrete points (temperature, time). Via ν these are converted to $(xpos, ypos)$ and pulled over discrete S . These values are then used to parameterize ρ which returns a color based on the parameters $(xpos, ypos)$ and position α, β on S_k that ρ is evaluated on.

The scatter plot in figure 12 can be defined as $Q(xpos, ypos)(\alpha, \beta)$ where color $\rho_{RGB} = (0, 0, 0)$ is defined as part of Q and $s = (\alpha, \beta)$ defines the region on S . The position of

this swatch of color can be computed relative to the location on the disc S_k as shown in figure 12:

$$x = size * \alpha \cos(\beta) + xpos \quad (31)$$

$$y = size * \alpha \sin(\beta) + ypos \quad (32)$$

221 such that $\rho(s) = (x, y, 0, 0, 0)$ colors the point (x,y) black. Here *size* can either be defined
 222 inside Q or it could also be a parameter in V that is passed along with (xpos, ypos). As
 223 seen in figure 12, a scatter has a direct mapping from a region on S_k to its corresponding k .




Figure 13: The line fiber $(time, temp)$ is thickened with the derivative $(time', temperature')$ because that information will be necessary to figure out the tangent to the point to draw a line. This is because the line needs to be pushed perpendicular to the tangent of (xpos, ypos). The data is converted to visual characteristics (xpos, ypos). The α coordinates on S specifies the position of the line, the β coordinate specifies thickness.

In contrast to the scatter, the line plot $Q(xpos, \hat{n}_1, ypos, \hat{n}_2)(\alpha, \beta)$ shown in fig 13 has a ξ function that is not only parameterized on k but also on the α distance along k and corresponding region in S . The line also exemplifies the need for the jet since the line needs to know the tangent of the data to draw an envelope above and below each (xpos,ypos) such that the line appears to have a thickness. The magnitude of the slope is

$$|n| = \sqrt{n_1^2 + n_2^2} \quad (33)$$

such that the normal is

$$\hat{n}_1 = \frac{n_1}{|n|}, \hat{n}_2 = \frac{n_2}{|n|} \quad (34)$$

which yields components of ρ

$$x = xpos(\xi(\alpha)) + width * \beta \hat{n}_1(\xi(\alpha)) \quad (35)$$

$$y = ypos(\xi(\alpha)) + width * \beta \hat{n}_2(\xi(\alpha)) \quad (36)$$

224 where (x,y) look up the position $\xi(\alpha)$ on the data. At that point, we also look up the the
 225 derivatives \hat{n}_1, \hat{n}_2 which are then multiplied by a *width* parameter to specify the thickness.
 226 As with the *size* parameter in scatter, *width* can be defined in Q or as a component of V .




Figure 14: The only visual parameter an image requires is color since ξ encodes the mapping between position in data and position in graphic.

The image $Q(xpos, ypos, color)$ in figure 14 is a direct lookup into $\xi : S \rightarrow K$. Since K is 2D continuous space, the indexing variables (α, β) define the distance along the space. This is then used by ξ to map into K to lookup the color values

$$R = R(\xi(\alpha, \beta)) \quad (37)$$

$$G = G(\xi(\alpha, \beta)) \quad (38)$$

$$B = B(\xi(\alpha, \beta)) \quad (39)$$

227 that the data values have been mapped into. In the case of an image, the indexing mapper
 228 ξ may do some translating to a convention expected by Q , for example reorientng the array
 229 such that the first row in the data is at the bottom of the graphic.

230 1.3.5 Assembly factory \hat{Q}

231 The graphic base space S is not accessible in many architectures, including Matplotlib;
 232 instead we can construct a factory function \hat{Q} over K that can build a Q . As shown in
 233 eq 18, Q is a bundle map $Q : \xi^*V \rightarrow H$ where ξ^*V and H are both bundles over S .




Figure 15: The pullback of the visual bundle ξ^*V is the replication of a μ over all points s that map back to a single k . Because the μ is the same, we can construct a \hat{Q} on μ over k that will fabricate the Q for the equivalent region of s associated to that k

The preimage of the continuity map $\xi^{-1}(k) \subset S$ is such that many graphic continuity points $s \in S_K$ go to one data continuity point k ; therefore, by definition the pull back of μ

$$\xi^*V|_{\xi^{-1}(k)} = \xi^{-1}(k) \times P \quad (40)$$

copies the visual fiber P over the the points s in graphic space S that correspond to one k in data space K . This set of points s are the preimage $\xi^{-1}(k)$ of k .

This copying is illustrated in figure 15, where the 1D fiber $P \hookrightarrow V$ over K is copied repeatedly to become the 2D fiber $P^*\mu \hookrightarrow \xi^*V$ with identical components over S . Given the section $\xi^*\mu$ pulled back from μ and the point $s \in \xi^{-1}(k)$, there is a direct map from μ on a point k , there is a direct map from the visual section over data base space $(k, \mu(k)) \mapsto (s, \xi^*\mu(s))$ to the visual section $\xi^*\mu$ over graphic base space. This map means that the pulled back section $\xi^*\mu(s) = \xi^*(\mu(k))$ is the section μ copied over all s . This means that $\xi^*\mu$ is identical for all s where $\xi(s) = k$, which is illustrated in figure 15 as each dot on P is equivalent to the line intersection $P^*\mu$.

Given the equivalence between μ and $\xi^*\mu$ defined above, the reliance on S can be factored out. When Q maps visual sections into graphics $Q : \Gamma(\xi^*V) \rightarrow \Gamma(H)$, if we restrict Q input to the pulled back visual section $\xi^*\mu$ then

$$\rho(s) := Q(\xi^*\mu)(s) \quad (41)$$

the graphic section ρ evaluated on a visual region s is defined as the assembly function Q with input pulled back visual section $\xi^*\mu$ also evaluated on s . Since the pulled back visual section $\xi^*\mu$ is the visual section μ copied over every graphic region $s \in \xi^{-1}(k)$, we can define a Q factory function

$$\hat{Q}(\mu(k))(s) := Q((\xi^*\mu)(s)) \quad (42)$$

where the assembly function \hat{Q} that takes as input the visual section on data μ is defined to be the assembly function Q that takes as input the copied section $\xi^*\mu$ such that both functions are evaluated over the same location $\xi^{-1}(k) = s$ in the base space S .

247 Factoring out s from equation 42 yields $\hat{Q}(\mu(k)) = Q(\xi^* \mu)$ where Q is no longer bound
 248 to input but \hat{Q} is still defined in terms of K . In fact, \hat{Q} is a map from visual space to
 249 graphic space $\hat{Q} : \Gamma(V) \rightarrow \Gamma(H)$ locally over k such that it can be evaluated on a single
 250 visual record $\hat{Q} : \Gamma(V_k) \rightarrow \Gamma(H|_{\xi^{-1}(k)})$. This allows us to construct a \hat{Q} that only depends
 251 on K , such that for each $\mu(k)$ there is part of $\rho|_{\xi^{-1}(k)}$. The construction of \hat{Q} allows us
 252 to retain the functional map reduce benefits of Q without having to majorly restructure
 253 the existing pipeline for libraries that delegate the construction of ρ to a back end such as
 254 Matplotlib.

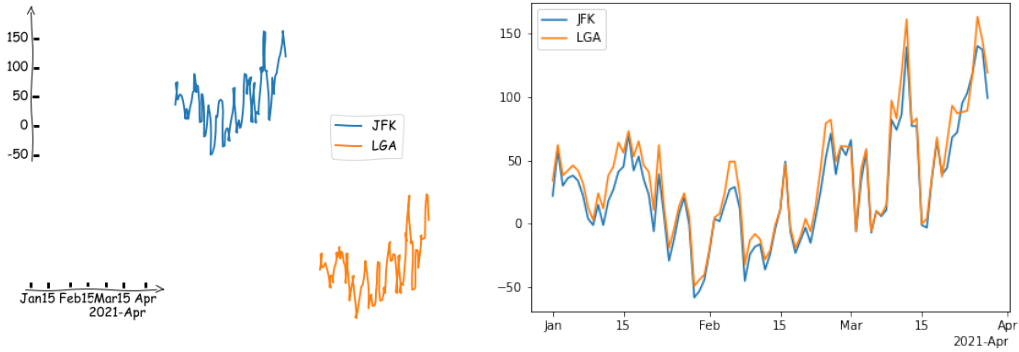
255 1.3.6 Sheafs

The restriction maps of a sheaf describe how local τ can be glued into larger sections [34, 35]. As part of the definition of local triviality, there is an open neighborhood $U \subset K$ for every $k \in K$. We can define the inclusion map $\iota : U \rightarrow K$ which pulls E over U

$$\begin{array}{ccc} \iota^* E & \xrightarrow{\iota^*} & E \\ \pi \downarrow \uparrow \iota^* \tau & & \pi \downarrow \uparrow \tau \\ U & \xrightarrow{\iota} & K \end{array} \quad (43)$$

256 such that the pulled back $\iota^* \tau$ only contains records over $U \subset K$. By gluing $\iota^* \tau$ together, the
 257 sheaf is putting a continuous structure on local sections which allows for defining a section
 258 over a subset in K . That section over subset K maps to the graphic generated by A for
 259 visualizations such as sliding windows[36, 37] streaming data, or navigation techniques such
 260 as pan and zoom[38].

261 1.3.7 Composition of Artists: +



(a) Each component of a visualization is generated via a different artist.
 (b) This visualization is the disjoint union of the artists in Figure 16a.

Figure 16

262 To build graphics that are the composites of multiple artists, we define a simple addition
 263 operator that is the disjoint union of fiber bundles E . For example, a scatter plot E_1 and a

264 line plot E_2 have different K that are mapped to separate S . To fully display both graphics,
 265 the composite graphic $A_1 + A_2$ needs to include all records on both K_1 and K_2 , which are the
 266 sections on the disjoint union $K_1 \sqcup K_2$. This in turn yields disjoint graphics $S_1 \sqcup S_2$ rendered
 267 to the same image. Constraints can be placed on the disjoint union such as that the fiber
 268 components need to have the same ν position encodings or that the position μ need to be in a
 269 specified range. There is a second type of composition where E_1 and E_2 share a base space
 270 $K_2 \hookrightarrow K_1$ such that the the artists can be considered to be acting on different components
 271 of the same section. This type of composition is important for creating visualizations where
 272 elements need to update together in a consistent way, such as multiple views [39, 40] and
 273 brush-linked views[41, 42].

274 **1.3.8 Equivalence class of artists A'**

275 It is impractical to implement an artist for every single graphic; instead we implement an
 276 approximation of an the equivalence class of artists $\{A \in A' : A_1 \equiv A_2\}$. Roughly, equivalent
 277 artists have the same fiber bundle V and same assembly function Q but act on different
 278 sections μ , but we will formalize the definition of the equivalence class in future work. As
 279 a first pass for implementation purposes, we identify a minimal P associated with each A'
 280 that defines what visual characteristics of the graphic must originate in the data such that
 281 the graphic is identifiable as a given chart type.

282 For example, a scatter plot of red circles is the output of one artist, a scatter plot of green
 283 squares the output of another. These two artists are equivalent since their only difference
 284 is in the literal visual encodings (color, shape). Shape and color could also be defined in Q
 285 but the position must come from the fiber $P = (xpos, ypos)$ since fundamentally a scatter
 286 plot is the plotting of one position against another[43]. We also use this criteria to identify
 287 derivative types, for example the bubble chart[44] is a type of scatter where by definition the
 288 glyph size is mapped from the data. The criteria for equivalence class membership serves
 289 as the basis for evaluating invariance[kindlmann2014algebraic].